

DATA EXPLORATION

Working with datasets and producing descriptive statistics

A few basic things first...

Working directory

- The working directory is the directory in which you are working in.
- To check your current working directory, use the function `getwd()`.
- To change your working directory, use the function `setwd(directory_path)`.
- Tips: running “`setwd(choose.dir())`” allows you to select the directory interactively.

Using packages

- Packages are a collection of codes and functions written by a third party and is not included in the base R.
- To install and download a package, use the function `install.packages("package_name")`.
- Then to load the package, use the function `library(package_name)`.

Using functions

- ❑ For built-in functions, using `help(function_name)` will give the documentation of the function which includes:
 - ❑ Description of the function
 - ❑ Arguments and its default value
 - ❑ Values the function returns
- ❑ The argument of a function takes the default value if it is not specified.
- ❑ When running a function, you can specify the argument by writing `fn(argument_name=value)`, or just setting `f(value)`.
- ❑ If you do not specify the argument name, then R will specify the argument by position.

Using functions

- Suppose we have:

```
fn_name(arg1 = value1, arg2 = value2)
```

- ▣ The name of the function is `fn_name`.
- ▣ Two arguments: `arg1` and `arg2`.
- ▣ Default value for `arg1` is `value1`, default value for `arg2` is `value2`.
- ▣ Running `fn_name()` is equivalent to running `fn_name(arg1=value1, arg2=value2)`.
- ▣ Running `fn_name(a)` is equivalent to `fn_name(arg1=a, arg2=value2)`.
- ▣ Running `fn_name(a, b)` is equivalent to `fn_name(arg1=a, arg2=b)`.
- ▣ Running `fn_name(arg2=a, arg1=b)` is equivalent to `fn_name(arg1=b, arg2=a)` or `fn_name(b, a)`.

Datasets in R

Data frames

- A data frame is a table or a two-dimensional array-like structure in which each column contains values of one variable and each row contains one set of values from each column.
- What this basically mean:
 - ▣ Data frames are like a combination of matrix and list (which we will explore later), each column corresponds to a variable and each row corresponds to a sample/observation.
 - ▣ One sample or observations has multiple variables.
 - ▣ E.g.:
 - Row corresponds to each student.
 - Column 1 is students' assignment marks.
 - Column 2 is students' test marks.
 - ▣ Like matrix: has columns and rows.
 - But unlike matrix, each columns can be of different class.
 - ▣ Like list: has variable names and can call them.

Creating data frame

- Use `data.frame()` function.
- Example:

```
> x <- 1:10
> y <- x^2
> dat <- data.frame(col1=x, col2=y)
> print(dat)
   col1 col2
1     1    1
2     2    4
3     3    9
4     4   16
... 
```

- You can also convert a matrix to a data frame using `as.data.frame()`.

Pulling out contents of a data frame

- We use dollar sign `$` to extract the variable nested inside of a data frame.
- Example:

```
> x <- 1:10
> y <- x^2
> dat <- data.frame(col1=x, col2=y)
> print(dat$col1)
[1] 1 2 3 4 5 6 7 8 9 10
> print(dat$col2)
[1] 1 4 9 16 25 36 49 64 81 100
```

Selecting a subset of observations with some criteria

- If you want to select/print observations with some conditions/criteria, you can use the `[]` notation like we did for matrix.
- Alternatively, you can use the `subset()` function.

```
subset(data_frame_name, condition)
```

Selecting a subset of observations with some criteria

□ Example:

```
> dat[1:5, ]  
  col1 col2  
1     1     1  
2     2     4  
3     3     9  
4     4    16  
5     5    25  
> dat[,1]  
[1]  1  2  3  4  5  6  7  8  9 10
```

Selecting a subset of observations with some criteria

□ Example:

```
> dat[dat$col2<=25, ]  
  col1 col2  
1     1     1  
2     2     4  
3     3     9  
4     4    16  
5     5    25  
> subset(dat,col2<=20)  
  col1 col2  
1     1     1  
2     2     4  
3     3     9  
4     4    16
```

Some functions for data frames

str() function

- The function `str()` can be used to quickly look into the structure of an object, including a data frame.
- It will give the number of observation as well as variable's names and classes
- Example:

```
> str(dat)
'data.frame':      10 obs. of  2 variables:
 $ coll: int  1  2  3  4  5  6  7  8  9 10
 $ col2: num  1  4  9 16 25 36 49 64 81 100
```

names() function

- The function `names()` can be used to list down all variable names (or column names) in a data frame.
- It can also be used to modify the column names
- Example:

```
> names(dat)
[1] "col1" "col2"
> names(dat) <- c("x", "y") #change the var names
> names(dat)
[1] "x" "y"
```


Row and column names

- You can also use `colnames()` to get the column names, and `rownames()` to get the row names of a data frame/matrix.
- Like the `names()` function, you can also use these to modify the column or row names.

Row and column names

□ Example:

```
> colnames(dat)
[1] "x" "y"
> rownames(dat)
[1] "1" "2" "3" "4" "5" "6" "7" "8" "9" "10"
> rownames(dat)[1] <- "No.1"
> dat
```

	x	y
No.1	1	1
2	2	4
3	3	9
4	4	16
5	5	25
6	6	36
7	7	49
8	8	64
9	9	81
10	10	100

```
> rownames(dat)[1] <- "1"
```

View() and fix() functions

- You can use `View()` function (capital V) to view the whole data frame on a separate window. But you can't edit it.
- If you want to edit it interactively, you can use `fix()` function.
- Example:

```
> View(dat)  
> fix(dat)
```

head() and tail() functions

- Sometimes you don't want to view all the dataset, but only a few just to see what it looks like.
- You can use the `head()` or `tail()` functions to print out the first or the last n rows in the dataset.
- The default value for n is 6, but you can modify it by specifying `n=value` in the argument of the functions.

head() and tail() functions

□ Example:

```
> head(dat)
```

	x	y
1	1	1
2	2	4
3	3	9
4	4	16
5	5	25
6	6	36

```
> tail(dat)
```

	x	y
5	5	25
6	6	36
7	7	49
8	8	64
9	9	81
10	10	100

More functions for data frames

- `data()`:
 - ▣ R has built-in datasets. To load these datasets, use `data(dataset_name)`
 - ▣ E.g.: `data(mtcars)`

- `na.omit()`:
 - ▣ Remove rows with missing values (coded as NA in the dataset).
 - ▣ Useful for analysis if you want to remove them.

More functions for data frames

- `merge()`:
 - ▣ Merge two data frames by common columns or row names.
 - ▣ In some ways similar to `rbind()` and `cbind()` but better as it merge based on column or row names.

- `lapply()` / `sapply`:
 - ▣ Apply function to each variables in the data frame.
 - ▣ `lapply()` returns a list. `sapply()` returns a vector.
 - ▣ E.g.:
 - `lapply(dat, FUN=mean)`
 - `sapply(dat, FUN=mean)`

More functions for data frames

□ `which()`:

- ▣ Used to determine which rows satisfy a condition.
- ▣ E.g.:

```
> which(dat$y > 60)
[1]  8  9 10
> which(dat$x < 5)
[1] 1 2 3 4
```

□ `attach()`:

- ▣ Attach a data frame to R search path.
- ▣ When a data frame is attached, there is no need to use `$` to call the variable.
- ▣ The function `detach()` must be used to detach the data frame.

Importing and exporting data

Importing data into R

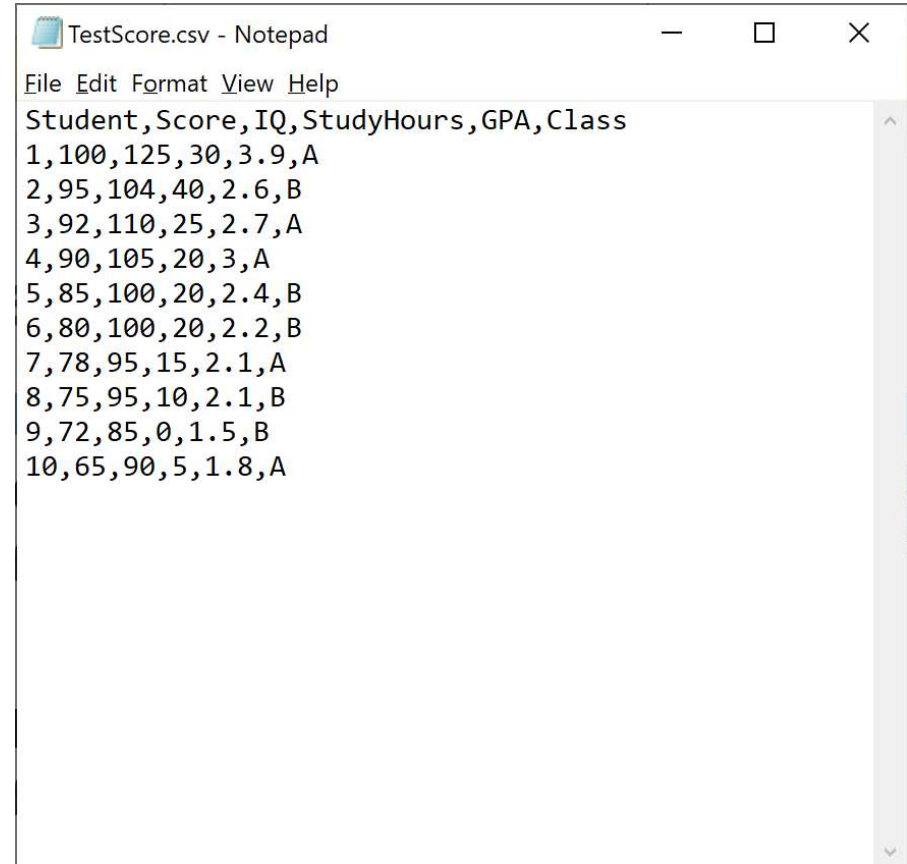
- We can type them down manually and use the `data.frame()`.
- For large amount of data, this is inefficient.
- It might be better to have the data in a separate txt, csv or excel file.

Reading from CSV or TXT file

- Use `read.table(...)` function.
- Argument: `read.table(file_location, header, sep)`
 - ▣ `file_location`: The location of the txt or csv. File location uses “/” , not “\”.
 - ▣ `header`: TRUE if there is a header (column name) in the file. Otherwise, FALSE.
 - ▣ `sep`: The separator between items in the file.

Reading from CSV or TXT file

- ❑ Open the file with notepad first to see the structure.
- ❑ In this case, we have header for each column.
 - ▣ `header=TRUE`
- ❑ And columns are separated with a comma “,”.
 - ▣ `sep=","`



```
TestScore.csv - Notepad
File Edit Format View Help
Student,Score,IQ,StudyHours,GPA,Class
1,100,125,30,3.9,A
2,95,104,40,2.6,B
3,92,110,25,2.7,A
4,90,105,20,3,A
5,85,100,20,2.4,B
6,80,100,20,2.2,B
7,78,95,15,2.1,A
8,75,95,10,2.1,B
9,72,85,0,1.5,B
10,65,90,5,1.8,A
```

Reading from CSV or TXT file

□ Example:

```
> TestScore <- read.table("C:/Users/Hilmi  
Majid/OneDrive/PnP/2021 Sem  
1/STQD6214/datasets/TestScore.csv",  
header=TRUE, sep=",")  
> head(TestScore)  
  Student Score  IQ StudyHours GPA Class  
1        1   100 125         30 3.9     A  
2        2    95 104         40 2.6     B  
3        3    92 110         25 2.7     A  
4        4    90 105         20 3.0     A  
5        5    85 100         20 2.4     B  
6        6    80 100         20 2.2     B  
> names(TestScore)  
[1] "Student"      "Score"        "IQ"  
[4] "StudyHours"   "GPA"          "Class"
```

Tips and tricks

- You can choose the file interactively using `file.choose()`.
 - ▣ Eg: `TestScore <- read.table(file.choose(), header=TRUE, sep=",")`
- For CSV file, an easier function to use is `read.csv()` function.
 - ▣ Eg: `TestScore <- read.csv(file.choose())`
- If your CSV file is large, you can use `read_csv()` function in `readr` package for faster loading time.
- Lastly, you can use RStudio to import data interactively.

Reading from MS Excel file

- Reading from Excel file (xls or xlsx) requires `readxl` package:
- Example:

```
> install.packages("readxl")  
> library(readxl)  
> oxygen <- read_excel("C:/Users/Hilmi  
Majid/OneDrive/PnP/1920 Sem 1/STQS3113/Slides/R  
tutorial/oxygen.xlsx")
```

Exporting csv file from R

- To save a data frame into a CSV file, use `write.csv()` function

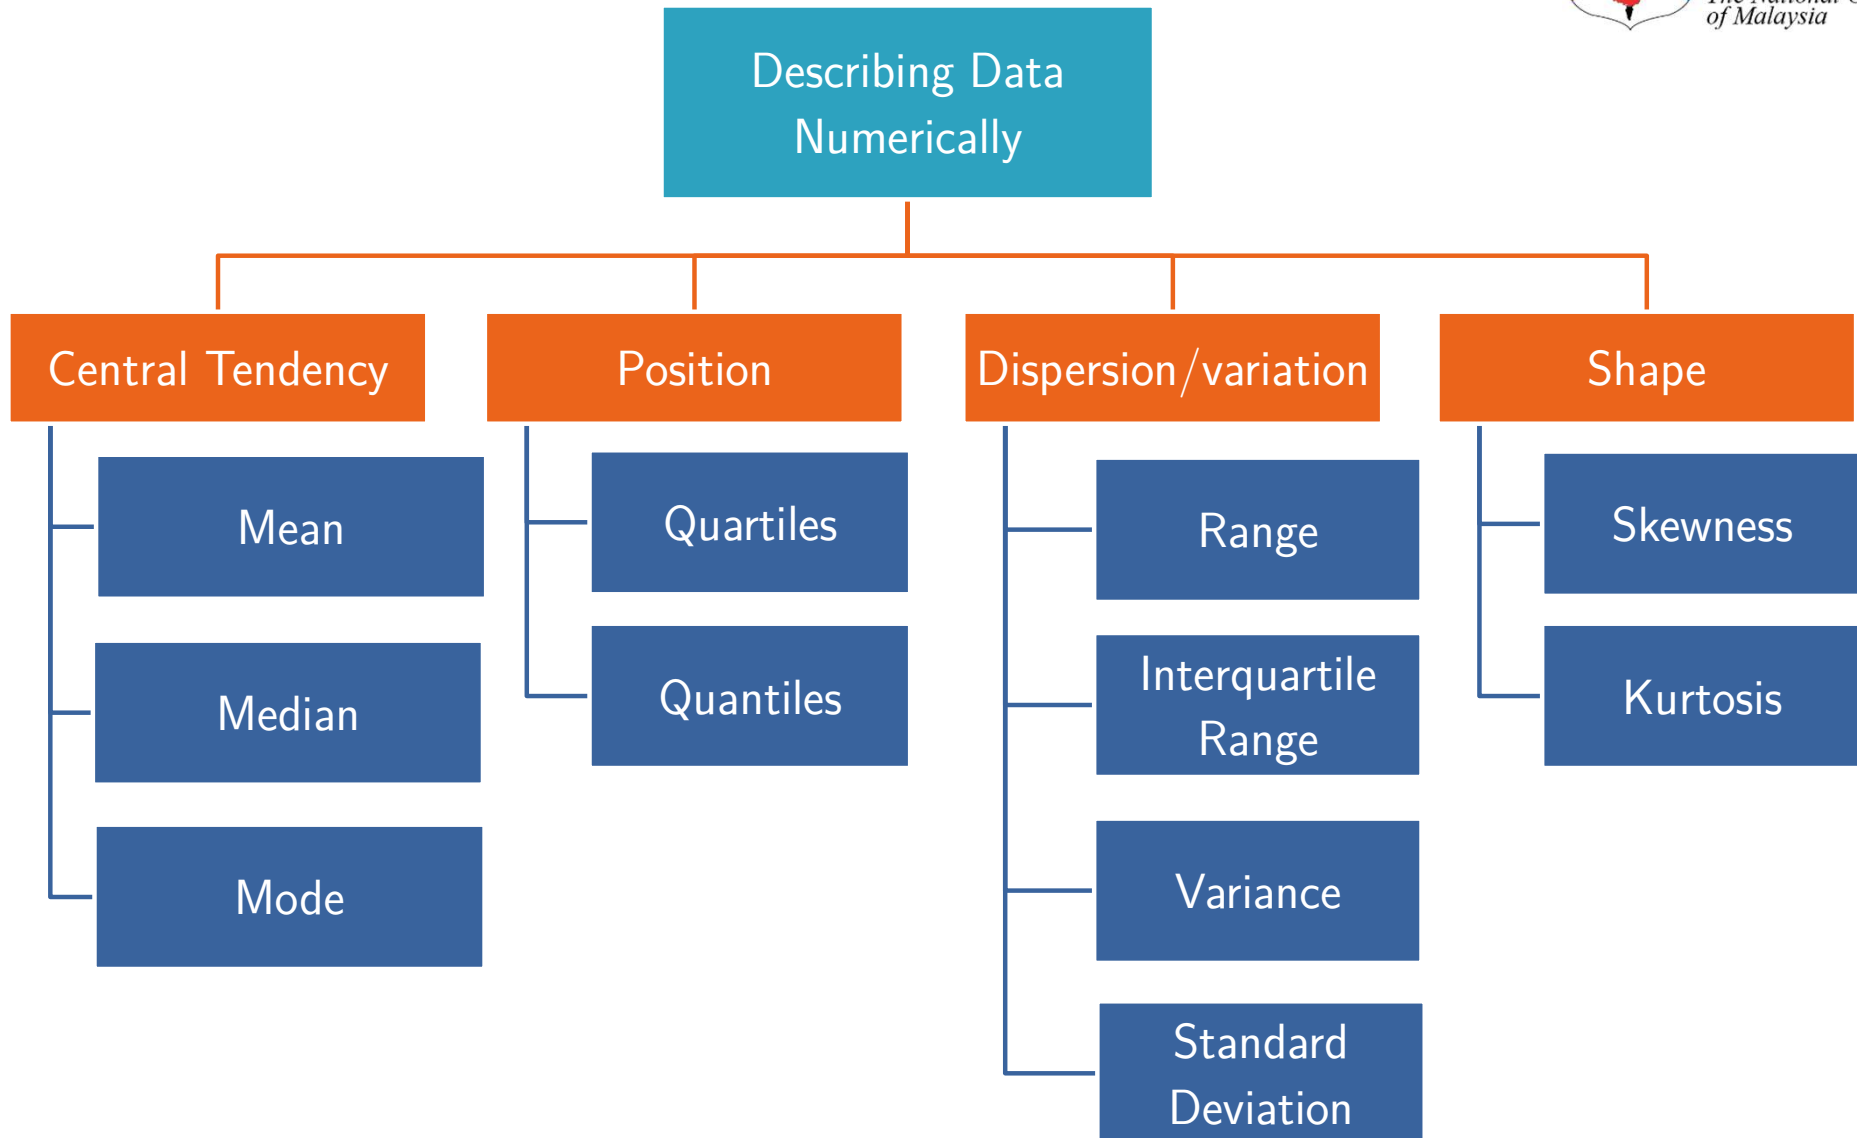
- Example:

```
write.csv(TestScore, file="TestScore.csv",  
          row.names=FALSE)
```

- If the data frame is too large, you can use `write_csv()` function from the `readr` package.
- (If you want to save any object in R, use the `save()` function)

Numerical descriptive statistics

Mean, median, mode, variance, standard deviation, etc



Summary statistics

- The `summary()` function when applied to vector or data frame gives the mean, min, max, first quarter, median, third quarter of each variables.
- If we want more summary statistics, we can use `describe()` function in the `psych` package.

Summary statistics

□ Example:

```
> summary(TestScore)

      Student          Score          IQ
Min.      : 1.00    Min.      : 65.00    Min.      : 85.0
1st Qu.: 3.25    1st Qu.: 75.75    1st Qu.: 95.0
Median : 5.50    Median : 82.50    Median :100.0
Mean     : 5.50    Mean      : 83.20    Mean      :100.9
3rd Qu.: 7.75    3rd Qu.: 91.50    3rd Qu.:104.8
Max.     :10.00    Max.      :100.00    Max.      :125.0

      StudyHours          GPA          Class
Min.      : 0.00    Min.      :1.500    A:5
1st Qu.:11.25    1st Qu.:2.100    B:5
Median :20.00    Median :2.300
Mean      :18.50    Mean      :2.430
3rd Qu.:23.75    3rd Qu.:2.675
Max.      :40.00    Max.      :3.900
```

Summary statistics

□ Example:

```
> install.packages("psych")
> library(psych)
> describe(TestScore)
```

	vars	n	mean	sd	median	trimmed	mad	min
Student	1	10	5.50	3.03	5.5	5.50	3.71	1.0
Score	2	10	83.20	11.10	82.5	83.38	12.60	65.0
IQ	3	10	100.90	11.22	100.0	99.88	7.41	85.0
StudyHours	4	10	18.50	11.80	20.0	18.12	11.12	0.0
GPA	5	10	2.43	0.68	2.3	2.36	0.52	1.5
Class*	6	10	1.50	0.53	1.5	1.50	0.74	1.0

	max	range	skew	kurtosis	se
Student	10.0	9.0	0.00	-1.56	0.96
Score	100.0	35.0	-0.06	-1.43	3.51
IQ	125.0	40.0	0.65	-0.32	3.55
StudyHours	40.0	40.0	0.14	-0.99	3.73
GPA	3.9	2.4	0.73	-0.27	0.21
Class*	2.0	1.0	0.00	-2.19	0.17

Table of frequency

- If the vector or variable is discrete, use `table()` function to create a table for the frequency of each value.

- Example:

```
> x <- c(3,2,2,4,2,1,2,2,3,4)
> table(x)
x
1 2 3 4
1 5 2 2
```


Measures of central tendency

Mean, median, mode

Mean

- What it is:
 - ▣ The sum of all the data entries divided by the number of entries.

$$\text{Mean, } \bar{x} = \frac{\sum x}{n}$$




Mean

- Use `mean()` function to calculate the mean of a vector.
- Example:

```
> mean(TestScore$Score)
[1] 83.2
> mean(TestScore$IQ)
[1] 100.9
> mean(TestScore$StudyHours)
[1] 18.5
```

Median

- What it is:
 - ▣ The value that lies in the middle of the data when the data set is ordered.



Median

□ Use `median()` function to calculate the median of a vector.

□ Example:

```
> median(TestScore$Score)
[1] 82.5
> median(TestScore$IQ)
[1] 100
> median(TestScore$StudyHours)
[1] 20
```

Mode

- What it is:
 - ▣ The data entry that occurs with the highest number of frequency.
- If no entry is repeated the data set has no mode.
- If two entries occur with the same greatest frequency, each entry is a mode (bimodal).

Mode

- If the vector or variable is discrete, we can use the `table()` function to find the frequency of each value and find the one with the highest frequency.
- Example:

```
> x <- c(3,2,2,4,2,1,2,2,3,4)
> table(x)
x
1 2 3 4
1 5 2 2
> names(table(x))[which.max(table(x))]
[1] "2"
```

Measures of position

Quartiles, quantiles

Quartile

- What it is:
 - ▣ Quartiles divide the distribution into four equal groups.
- The boundaries for the groups are denoted by Q_1 , Q_2 , Q_3 .



- Eg: Approximately 75% of values in a ranked data set are more than Q_1 .
- Note that the median is equal to Q_2 .

Quantile

- The `quantile()` function by default gives the min, Q_1 , Q_2 , Q_3 , and the max of a vector.
- But the argument can be modified to find any quantiles.
- Example:

```
> quantile(TestScore$Score)
      0%      25%      50%      75%     100%
 65.00  75.75  82.50  91.50 100.00
> quantile(TestScore$Score, probs=c(0.1,0.9))
 10%   90%
71.3  95.5
```


Measures of dispersion

Variance, standard deviation, range, interquartile range

Variance & standard deviation

- What it is:
 - ▣ The values tell how closely the observations are to the mean.
 - ▣ Variance and standard deviation are the most used measures of dispersion.

- Sample variance and sample standard deviation:

$$s^2 = \frac{\sum (x - \bar{x})^2}{n - 1}, \quad s = \sqrt{\frac{\sum (x - \bar{x})^2}{n - 1}}$$

- The lower the value, the closer the observations are to the sample mean.

Variance & standard deviation

□ The `var()` and `sd()` functions calculate these values.

□ Example:

```
> var(TestScore$Score)
[1] 123.2889
> sd(TestScore$Score)
[1] 11.10355
```

□ Note that these are sample variance and standard deviation, not population variance and standard deviation.

Range

- What it is: $\text{Range} = \text{Largest value} - \text{smallest value}$
- We can use `max()` and `min()` to get the largest and smallest values respectively to calculate range.
- The function `range()` also gives the largest and smallest value.
- Example:

```
> max(TestScore$Score) - min(TestScore$Score)
[1] 35
> range(TestScore$Score)
[1] 65 100
```

Interquartile range

- What it is:
 - ▣ The difference between Q_1 and Q_3 .
 - ▣ $IQR = Q_3 - Q_1$
- The `IQR()` function calculates the interquartile range.
- Example:

```
> quantile(TestScore$Score, c(0.25,0.75))  
 25%    75%  
75.75 91.50  
> IQR(TestScore$Score)  
[1] 15.75
```

Shape of the distribution

Skewness and kurtosis


Skewness

- Skewness measures the degree of asymmetry exhibited by the data.

$$\text{skewness} = \frac{\sum_{i=1}^n (x_i - \bar{x})^3}{ns^3}$$

- Skewness:
 - ▣ zero: data is symmetric about the mean
 - ▣ negative: data is skewed to the left
 - ▣ positive: data is skewed to the right

Skewness



- We can use `skewness()` function from `e1071` package to calculate this value.


- Example:

```
> library(e1071)
> skewness(TestScore$Score)
[1] -0.05638212
```


Kurtosis

- Kurtosis is a measure of the heaviness of the tail of a distribution.
- It is often used to represent how often outliers occur
- Excess kurtosis is the tailedness of a distribution relative to a normal distribution.
- Excess kurtosis:
 - ▣ High/positive (leptokurtic): fat tails
 - ▣ Low/negative (platykurtic): thin tails
 - ▣ Medium/zero (mesokurtic): medium tails

Kurtosis



□ Example:

```
> library(e1071)
> kurtosis(TestScore$Score)
[1] -1.434192
```

Functions & descriptions

Function	Description
<code>data.frame()</code>	Creates a new data frame class object.
<code>str()</code>	Gives the structure of an object.
<code>subset()</code>	Creates a subset of the data frame based on the specified condition.
<code>names()</code>	Gives the available variable names in the data frame/list.
<code>colnames()</code> , <code>rownames()</code>	Gives the column and row names in the data frame/matrix.
<code>head()</code> , <code>tail()</code>	Gives the first few rows and last few rows in the data frame.
<code>getwd()</code> , <code>setwd()</code>	Gives and sets the location of current working directory.
<code>read.table()</code>	Imports CSV or TXT files as data frame.
<code>read.csv()</code>	Imports CSV files as data frame.
<code>write.csv()</code>	Exports a data frame into a CSV file.

Functions & descriptions

Function	Description
<code>summary()</code>	When the argument/input is a vector, it gives some summary statistics including min, max, mean, and the quartiles.
<code>describe()</code>	Gives more summary statistics. This function requires <code>psych</code> package to be loaded first.
<code>table()</code>	Gives a table of frequency.
<code>mean(), median()</code>	Calculates the mean and median of a vector.
<code>quantile()</code>	Gives the sample quantiles. By default, it gives min, Q_1 , Q_2 , Q_3 and max.
<code>IQR()</code>	Calculates the interquartile range of a vector.
<code>var(), sd()</code>	Calculates the (sample) variance and standard deviation of a vector.
<code>skewness(), kurtosis()</code>	Gives skewness and kurtosis measures of a vector. These functions require <code>e1071</code> package to be loaded first.