

MONTE CARLO METHODS

Probability distributions in R, Monte Carlo simulations

Probability distributions in R

Probability distributions in R

- For most probability distributions, R has built-in functions to
 1. Calculate the pdf/pmf of a distribution.
 2. Calculate the cdf of the distribution.
 3. Calculate the quantile function of the distribution.
 4. Generate random numbers based on the distribution.

- All the functions related to above have the same “form” for all distributions.

Probability distributions in R

- For example, the functions related to a normal distribution are `dnorm()`, `pnorm()`, `qnorm()` and `rnorm()`.
- Functions related to a binomial distribution are `dbinom()`, `pbinom()`, `qbinom()` and `rbinom()`.
- Note:
 - ▣ “d” is for density. Used to calculate pdf/pmf.
 - ▣ “p” is for probability. Used to calculate cdf.
 - ▣ “q” is for quantile. Used to calculate quantile function given probability
 - ▣ “r” is for random. Used to generate random samples from the distribution.

Normal distribution

- Calculate pdf:

- ▣ `dnorm(x, mean = 0, sd = 1, log = FALSE)`

- Calculate cdf:

- ▣ `pnorm(q, mean = 0, sd = 1, lower.tail = TRUE, log.p = FALSE)`

- Calculate quantile function:

- ▣ `qnorm(p, mean = 0, sd = 1, lower.tail = TRUE, log.p = FALSE)`

- Generate random numbers:

- ▣ `rnorm(n, mean = 0, sd = 1)`

Example

- Let X be a random variable distributed following a normal distribution with mean 3 and standard deviation 1.5. Calculate the following probabilities.
 - ▣ $P(X < 2.54)$
 - ▣ $P(X > 3.5)$
 - ▣ $P(2.54 \leq X < 3.5)$

- Find the value of x such that
 - ▣ $P(X \leq x) = 0.32$
 - ▣ $P(X > x) = 0.75$

Example


```
> pnorm(2.54, mean=3, sd=1.5)
[1] 0.3795486
> pnorm(3.5, mean=3, sd=1.5, lower.tail=FALSE)
[1] 0.3694413
> pnorm(3.5, mean=3, sd=1.5) - pnorm(2.54, mean=3,
    sd=1.5)
[1] 0.2510101
> qnorm(0.32, mean=3, sd=1.5)
[1] 2.298452
> qnorm(0.75, mean=3, sd=1.5, lower.tail=FALSE)
[1] 1.988265
```

Example

- Let X be a random variable distributed following a normal distribution with mean 3 and standard deviation 1.5.
 - ▣ Generate 100 samples of X .
 - ▣ Draw a histogram of the generated samples where the height of the histogram represent the probability of each bin.
 - ▣ Add a line showing the pdf of X onto the drawn histogram.

Example

```
> x <- rnorm(100, mean=3, sd=1.5)
> hist(x, freq=FALSE)
> curve(dnorm(x, mean=3, sd=1.5), from=-2, to=8,
        col="red", add=TRUE)
```



Common probability distributions in R

Distributions	Functions			
Beta	<code>pbeta</code>	<code>qbeta</code>	<code>dbeta</code>	<code>rbeta</code>
Binomial	<code>pbinom</code>	<code>qbinom</code>	<code>dbinom</code>	<code>rbinom</code>
Chi-Square	<code>pchisq</code>	<code>qchisq</code>	<code>dchisq</code>	<code>rchisq</code>
Exponential	<code>pexp</code>	<code>qexp</code>	<code>dexp</code>	<code>rexp</code>
F	<code>pf</code>	<code>qf</code>	<code>df</code>	<code>rf</code>
Gamma	<code>pgamma</code>	<code>qgamma</code>	<code>dgamma</code>	<code>rgamma</code>
Log Normal	<code>plnorm</code>	<code>qlnorm</code>	<code>dlnorm</code>	<code>rlnorm</code>
Normal	<code>pnorm</code>	<code>qnorm</code>	<code>dnorm</code>	<code>rnorm</code>
Poisson	<code>ppois</code>	<code>qpois</code>	<code>dpois</code>	<code>rpois</code>
Student t	<code>pt</code>	<code>qt</code>	<code>dt</code>	<code>rt</code>
Uniform	<code>punif</code>	<code>qunif</code>	<code>dunif</code>	<code>runif</code>

Sampling from a vector

- The `sample()` function can be used to randomly pick a sample (whether it be character, numbers, etc) from a given vector.
- Syntax: `sample(x, size, replace = FALSE, prob = NULL)`
 - ▣ `x`: a vector from which the elements are to choose from
 - ▣ `size`: how many samples we want
 - ▣ `replace`: should the sampling done with replacement? (can a sample be picked more than once?)
 - ▣ `prob`: a vector of the probability of picking each elements of `x`.

Example

```
> sample(c(1,2,3), 5, replace=TRUE)
[1] 3 1 2 3 1
> sample(c("A", "B", "C", "D"), 2, prob = c(0.1, 0.2,
0.5, 0.2))
[1] "B" "C"
```

- In the first example above, we want to get 5 values from the vector (1, 2, 3) where each element of the vector has the same probability of being picked. We set `replace=TRUE` because an element can be picked more than once.
- In the second example above, we want to get two values from the vector (A, B, C, D), where the probability of A being picked is 0.1, B is 0.2, C is 0.5, and D is 0.2. Since we do not specify `replace` argument, by default, `replace` is `FALSE` and a value cannot be picked more than once.

Random number generating process

- Note that the numbers generated in computer are not exactly random. They are “pseudorandom”.
- But they are good enough for most applications.
- You can set the “seed” of your randomly generated numbers using `set.seed()` function

Example

```
> set.seed(1)
> rnorm(10)
[1] -0.6264538  0.1836433 -0.8356286  1.5952808
[5]  0.3295078 -0.8204684  0.4874291  0.7383247
[9]  0.5757814 -0.3053884
> set.seed(1)
> rnorm(10)
[1] -0.6264538  0.1836433 -0.8356286  1.5952808
[5]  0.3295078 -0.8204684  0.4874291  0.7383247
[9]  0.5757814 -0.3053884
```

- Notice the randomly generated numbers are exactly the same.


Monte Carlo methods

What are Monte Carlo methods?

- Basically, solving problems by simulations.
 - ▣ Estimating solution to problems cannot be solved analytically.
 - ▣ Testing methods or algorithms on simulated data.
 - ▣ Mimics experimental lab.

- Why use simulations?
 - ▣ For problems that cannot be solved analytically, sometimes simulations give easier solution
 - ▣ Able to perform analysis on a “perfect” data.

Example 1: Estimating the value of π



- Area of a unit square = 1
- Area of the quadrant = $\frac{\pi}{4}$
- Proportion of the area of quadrant over the area of square = $\frac{\pi}{4}$

Example 1: Estimating the value of π

- Algorithm:
 - ▣ Simulate the X and Y values for the coordinates of points in the unit square ($X \sim \text{Unif}(0,1)$ and $Y \sim \text{Unif}(0,1)$).
 - ▣ Count the proportion of simulated data that is inside the quadrant.
 - ▣ Estimate the value of π by





$$\pi \approx 4 \times (\text{proportion data inside quadrant})$$

Example 1: Estimating the value of π

```
### Estimating pi
```

```
estimate_pi <- function(n){  
  x <- runif(n=n, min=0, max=1)  
  y <- runif(n=n, min=0, max=1)  
  number_inside <- sum(y <= sqrt(1-x^2))  
  proportion <- number_inside/n  
  pi_est <- 4*proportion  
  return(pi_est)  
}
```

```
> estimate_pi(100)  
[1] 3.2  
> estimate_pi(1000)  
[1] 3.16  
> estimate_pi(1e6)  
[1] 3.140148  
> pi  
[1] 3.141593
```



Example 2: Inverse transform sampling

- Suppose we have a random variable X with cdf $F(x)$.
- Then it can be shown that the random variable $U = F(X)$ is uniformly distributed between 0 and 1.
- Theorem: Let $U \sim \text{Unif}(0,1)$ and F be a cdf. Then $F^{-1}(U)$ has cdf F .
- We can use this property to randomly generate the variable X .

Example 2: Inverse transform sampling

- Example: Suppose

$$f(x) = 6x^2(1 - x^3), x \in [0,1]$$

$$F(x) = 1 - (1 - x^3)^2$$

- Algorithm:

- ▣ Randomly generate U from $\text{Unif}(0,1)$.

- ▣ Calculate $X = F^{-1}(U)$

- First, we have to calculate $F^{-1}(U)$.

$$F^{-1}(U) = \left(1 - \sqrt{1 - U}\right)^{\frac{1}{3}}$$

Example 2: Inverse transform sampling

```
### Inverse transform sampling


generate_x <- function(n){
  u <- runif(n, min=0, max=1)
  x <- (1-sqrt(1-u))^(1/3)
  return(x)
}

# Compare histogram and pdf
x <- generate_x(1000)
hist(x, freq=FALSE)
curve(6*x^2*(1-x^3), from=0, to=1, add=TRUE,
      col="red")


# Compare empirical cdf and cdf
plot(ecdf(x))
curve(1-(1-x^3)^2, from=0, to=1, add=TRUE, col="red",
      lwd=2)
```


Example 2: Inverse transform sampling

Histogram of x



ecdf(x)



Functions & descriptions

Function	Description
<code>dnorm()</code>	Used to calculate pdf of normal distribution.
<code>pnorm()</code>	Used to calculate cdf of normal distribution.
<code>qnorm()</code>	Used to calculate quantile function of normal distribution.
<code>rnorm()</code>	Used to generate random samples from the normal distribution.
<code>sample()</code>	Used to get samples from a given list

- Note: Other distributions are also available and follow similar notations for the functions – `d`, `p`, `q`, `r`